

BUILDING A NEURAL NETWORK WITH TENSORFLOW & KERAS

1. From Theory to Practical

In the previous session, we learned the **fundamentals of Neural Networks**:

Neuron → Weights → Bias → Activation → Loss → Backpropagation → Weight Update

Now, we will take those concepts and **build an actual Neural Network using Python**.

Our goal is to understand:

How do we convert Neural Network theory into actual code?

2. What is TensorFlow?

TensorFlow is an open-source Machine Learning and Deep Learning library developed by Google.

It provides tools to:

- Build Neural Networks
- Train Neural Networks
- Evaluate models
- Make predictions

Simple Understanding

Think of TensorFlow as a **tool/framework that helps us create and train Deep Learning models**.

We don't have to manually write all the mathematical operations for:

Weights → Forward Pass → Loss → Backpropagation → Weight Updates

TensorFlow handles these operations for us.

3. What is Keras?

Keras is a high-level API integrated with TensorFlow.

It makes building Neural Networks easier by providing simple functions for:

- Creating layers
- Building models
- Training models
- Evaluating models
- Making predictions

Simple Understanding

Think of it like this:

TensorFlow → Deep Learning engine

Keras → Easy interface for building Neural Networks

So, in our code:

TensorFlow + Keras → Build and Train Neural Network

4. What Problem Are We Solving?

Now we will use the **Iris Dataset**.

Our goal is to predict the **species of an Iris flower** based on its measurements.

There are 3 species:

Setosa

Versicolor

Virginica

5. Understanding the Dataset

Each flower has four measurements:

- Sepal Length
- Sepal Width
- Petal Length
- Petal Width

These are our **Input Features (X)**.

The flower species is our **Target (y)**.

Simple Flow

Flower Measurements

↓

Neural Network

↓

Flower Species

In Simple Terms

We give the model the measurements of a flower and ask:

"Which type of Iris flower is this?"

6. Import Required Libraries

We first import the tools required for our project.

import tensorflow as tf

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
import matplotlib.pyplot as plt
```

We don't need to memorize every import.

The important ones are:

TensorFlow/Keras → Build the Neural Network

Scikit-learn → Dataset and data preprocessing

Matplotlib → Visualize training results

7. Load the Iris Dataset

```
iris = load_iris()
```

```
X = iris.data
```

```
y = iris.target.reshape(-1, 1)
```

X

Contains the input features:

Sepal Length, Sepal Width, Petal Length, Petal Width

y

Contains the target class:

0, 1, 2

These numbers represent the three Iris species.

Remember

X = Input

y = Output / Target

8. One-Hot Encoding

Our target currently contains:

0, 1, 2

We convert these values into a format suitable for multi-class classification.

```
encoder = OneHotEncoder(sparse_output=False)
```

```
y_encoded = encoder.fit_transform(y)
```

Example

$0 \rightarrow [1, 0, 0]$

$1 \rightarrow [0, 1, 0]$

$2 \rightarrow [0, 0, 1]$

This is called **One-Hot Encoding**.

Why?

Because we have **three different classes**, and our output layer will produce a probability for each class.

9. Train-Test Split

Now we divide our data into training and testing data.

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y_encoded,  
    test_size=0.2,  
    random_state=42  
)
```

Here:

80% → Training Data

20% → Testing Data

Why?

The model learns from the training data.

Then we test it using data that it has not used for training.

Simple Flow

Dataset

↓

80% Training | 20% Testing

10. Building the Neural Network

Now we create our actual Neural Network.

```
model = Sequential([
    Dense(10, input_shape=(X.shape[1],), activation='relu'),
    Dense(8, activation='relu'),
    Dense(y_encoded.shape[1], activation='softmax')
])
```

This is the **most important part**.

Let's understand it layer by layer.

11. Sequential Model

```
model = Sequential([
```

Sequential means we are creating the Neural Network **layer by layer in a sequence**.

Our flow will be:

Input → Hidden Layer → Hidden Layer → Output

12. First Hidden Layer

```
Dense(10, input_shape=(X.shape[1],), activation='relu')
```

This creates a **Dense layer with 10 neurons**.

Our Iris dataset has:

4 input features

So the flow is:

4 Inputs → 10 Neurons

The activation function is:

ReLU

We learned ReLU in the previous session.

13. Second Hidden Layer

```
Dense(8, activation='relu')
```

This creates another Dense layer containing:

8 neurons

Again, we use:

ReLU

So now our network looks like:

4 Inputs → 10 Neurons → 8 Neurons

14. Output Layer

```
Dense(y_encoded.shape[1], activation='softmax')
```

We have:

3 Iris Classes

Therefore, the output layer contains:

3 neurons

The activation function is:

Softmax

Why Softmax?

Because this is a **multi-class classification problem**.

The three output neurons represent:

Setosa

Versicolor

Virginica

Softmax gives probability-like values for these classes.

Example

Setosa → 0.05

Versicolor → 0.90

Virginica → 0.05

The highest value is:

Versicolor

Therefore, the model predicts:

Versicolor

15. Our Complete Neural Network

The complete architecture is:

Input

4 Features



10 Neurons

ReLU



8 Neurons

ReLU



3 Neurons

Softmax



Iris Species

Or simply:

4 → 10 → 8 → 3

This is the actual Neural Network we are building.

16. Compile the Model

Before training, we need to tell the model **how it should learn**.

```
model.compile(  
    optimizer='adam',  
    loss='categorical_crossentropy',  
    metrics=['accuracy']  
)
```

There are three important things here:

Optimizer

Adam

The optimizer helps update the model's weights during training.

Remember:

Prediction → Loss → Weight Update

Loss Function

Categorical Cross-Entropy

We use this because our problem is:

Multi-Class Classification

with:

One-Hot Encoded Targets

Metric

Accuracy

Accuracy tells us how many predictions are correct.

Example:

If 90 out of 100 predictions are correct:

Accuracy = 90%

17. Train the Model

Now we train our Neural Network.

```
history = model.fit(  
    X_train,  
    y_train,  
    validation_split=0.2,  
    epochs=50,  
    batch_size=8,  
    verbose=1  
)
```

The most important function here is:

`model.fit()`

It means:

Train the Neural Network using the training data.

18. What is an Epoch?

epochs=50

An **epoch** means one complete pass through the training data.

Here:

50 Epochs = The model gets 50 opportunities to learn from the training data.

During training, the model repeatedly:

Makes Prediction → Calculates Loss → Updates Weights

19. What is Batch Size?

batch_size=8

Instead of processing the entire training dataset at once, the model processes:

8 samples at a time

Then it updates the weights and moves to the next batch.

Simple Idea

Training Data

↓

8 Samples

↓

Update Weights

↓

Next 8 Samples

↓

Update Weights

20. What is Validation Split?

validation_split=0.2

This keeps aside **20% of the training data** for validation.

Validation helps us monitor how the model performs on data that is not being directly used for updating the weights.

So we can monitor:

Training Performance

and

Validation Performance

during training.

21. What Happens During Training?

Remember our previous theory:

Input

↓

Forward Pass

↓

Prediction

↓

Calculate Loss

↓

Backpropagation

↓

Update Weights

↓

Repeat

All of this learning process happens when we call:

```
model.fit()
```

This is where our **theory becomes practical code**.

22. Evaluate the Model

After training, we test the model on the test data.

```
loss, accuracy = model.evaluate(X_test, y_test)
```

```
print(f"\nTest Loss: {loss:.4f}")
```

```
print(f"Test Accuracy: {accuracy:.4f}")
```

We use:

```
model.evaluate()
```

to check how well our trained model performs on **unseen test data**.

Example

If the output is:

Test Accuracy: 0.9333

That means:

Approximately 93.33% of the test predictions were correct.

23. Training History

During training, we stored:

```
history = model.fit(...)
```

The history object stores information about the training process.

For example:

Training Accuracy

Validation Accuracy

Training Loss

Validation Loss

We can visualize this information using graphs.

24. Accuracy Graph

```
plt.plot(history.history['accuracy'],  
         label='Train Accuracy')
```

```
plt.plot(history.history['val_accuracy'],  
         label='Validation Accuracy')
```

This graph shows how accuracy changes across epochs.

What do we want?

Generally, we want:

Accuracy ↑

as training progresses.

It helps us understand whether the model is learning.

25. Loss Graph

```
plt.plot(history.history['loss'],  
         label='Train Loss')
```

```
plt.plot(history.history['val_loss'],  
         label='Validation Loss')
```

This graph shows how the loss changes during training.

Generally, we want:

Loss ↓

as training progresses.

Remember

Accuracy → Higher is better

Loss → Lower is better

26. Complete Workflow

Our complete Deep Learning workflow is:

Load Dataset

↓

Prepare X and y

↓

One-Hot Encode Target

↓

Train-Test Split

↓

Build Neural Network

↓

Compile Model

↓

****Train using **fit()**

↓

****Evaluate using **evaluate()**



Visualize Accuracy & Loss

27. Connect Today's Code With Yesterday's Theory

Yesterday we learned:

Neuron

Weights

Bias

Activation

Loss

Backpropagation

Gradient Descent

Today we implemented those concepts using:

TensorFlow + Keras

We don't manually write the complete mathematics of neural-network training.

Instead, TensorFlow/Keras provides the tools to build and train the network.

QUICK REVISION

TensorFlow

A Deep Learning framework used to build and train models.

Keras

A high-level API that makes building Neural Networks easier.

Sequential

Builds a model layer by layer.

Dense

Represents a fully connected Neural Network layer.

ReLU

Activation function commonly used in hidden layers.

Softmax

Used in our output layer for multi-class classification.

One-Hot Encoding

Converts class labels into a binary vector representation.

Optimizer

Helps update the model's weights during training.

Loss Function

Measures how wrong the model's predictions are.

Epoch

One complete pass through the training data.

Batch Size

Number of samples processed before a weight update.

Fit

Trains the model.

Evaluate

Checks model performance on test data.

Accuracy

Measures the percentage of correct predictions.

KEY IDEA TO REMEMBER

Yesterday:

We learned HOW a Neural Network works.

Today:

We learned HOW to BUILD and TRAIN that Neural Network using TensorFlow and Keras.

Our Model:

4 Inputs → 10 Neurons → 8 Neurons → 3 Outputs

Complete Process:

Data → Preprocessing → Neural Network → Training → Evaluation → Results

Most Important Functions:

`model = Sequential(...)`

↓

`model.compile(...)`

↓

`model.fit(...)`

↓

`model.evaluate(...)`

Build → Compile → Train → Evaluate